Machine Learning Project
Classifying human activities
(Project #3)

Elia Perantoni

June 5, 2022

1 Introduction

We want to build a machine learning classifier on the *Human Activity Recognition with Smartphones* dataset, available at the following link.

The data has been gathered by having 30 subjects performing 6 different physical activities (laying, standing, sitting, walking, walking upstairs and walking downstairs) while a waist-mounted smartphone records accelerometer and gyroscope state.

Each sample originally consisted of 3-axis linear acceleration and 3-axis angular velocity but has been feature-engineered, by the authors, to a grand total of 561 features, obtained by filtering, mixing and matching the primitive ones.

The dataset is already split in 7352 training samples and 2947 testing samples.

2 Class Encoding

To make the rest of the code easier, we encode the labels (given as strings) to integers using the `ClassMapper` class.

```
# train_y is ["WALKING", "STANDING", "WALKING", ... ]
cm = ClassMapper(train_y)
train_y = cm.transform(train_y) # train_y is [0, 1, 0, ... ]
```

3 PCA

We start by performing *Principal Component Analysis* to reduce the number of features and make classification easier.

The `pca` function takes in a $n \times d$ matrix (n is the number of samples and d is the number of features), a `retain` coefficient that indicates how many axes the data should be projected on, and returns a Python function that performs the dimensionality reduction.

```
pca_reduce = pca(train_X, retain=0.3, verbose=True)
train_X = pca_reduce(train_X)
```

`retain=0.3` means that the samples will be projected on $0.3 \cdot d$ axes. In our case, the samples are scaled from 561 features down to 168.


Figure 1: PCA processed data projected on first two features

4 LDA

To make classification even easier, we apply *Linear Discriminant Analysis* to transform the data such that classes are tightly packed and distant from each other. Since we have $k = 6$ classes, we will be projecting on $k - 1 = 5$ axes.

The implementation is functionally the same as *PCA*: a function takes in some data and produces a function that performs the *LDA* dimensionality reduction. The only differences are that we no longer need a `retain` argument but now require the labels `train_y`.

```
lda_reduce = lda(train_X, train_y, verbose=True)
train_X = lda_reduce(train_X)
```
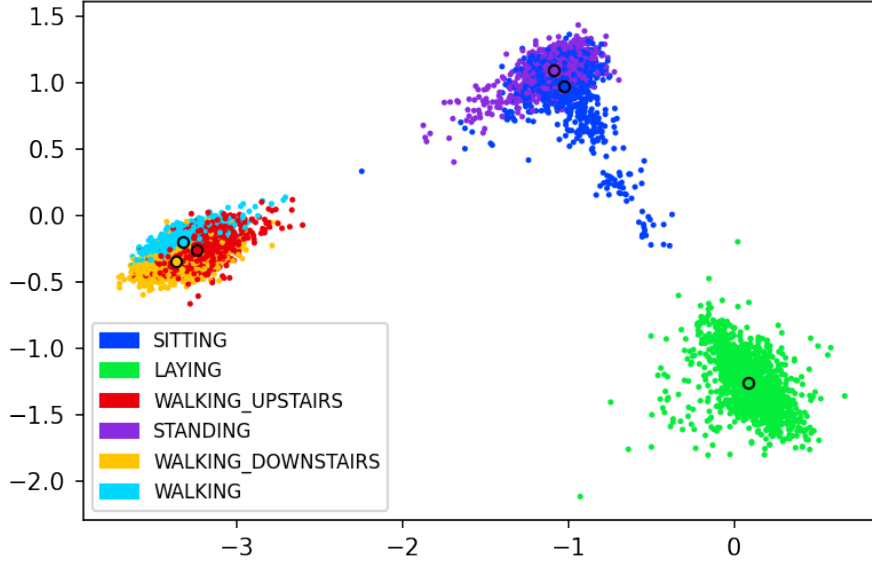

Figure 2: PCA+LDA processed data projected on first two features

5 Classification

Now that we have 7352 samples of 5 features each, we can start training a classification system.

I chose to go down the naive Bayesian route. There are $k = 6$ multivariate Gaussian distributions, one for each class, that attempt to enclose the region of $\mathbb{R}^5$ in which the samples of that class occur. The Gaussian distributions are trained simply by partitioning the training samples by class and, for each partition, computing the mean $\boldsymbol{\mu}$ and covariance matrix $\boldsymbol{\Sigma}$. Although the classes are quite evenly distributed, the prior probability $p(\omega_i) = \frac{|X \cap \omega_i|}{|X|}$ is also computed.

To make a new prediction on an unseen sample $\boldsymbol{x}$, simply take the class whose Gaussian distribution returns the greatest posterior probability.

$$\operatorname*{argmax}_{\omega_i} p(\omega_i \mid \boldsymbol{x}) = \frac{p(\boldsymbol{x} \mid \omega_i) \cdot p(\omega_i)}{p(\boldsymbol{x})} = p(\boldsymbol{x} \mid \omega_i) \cdot p(\omega_i)$$

The evidence probability $p(\boldsymbol{x})$ can be ignored because it only normalizes the posteriors, the winning class would still be the same.

6 Results

Our simple classifier achieves an accuracy of 0.9613. Compare this with other, more sophisticated classifiers (from the *SciKit-Learn* library).

Algorithm	Accuracy
Random Forest	0.9636
Naive Bayes *(ours)*	0.9613
SVM	0.9613
Multi-Layer Perceptron Neural Network	0.9613
K-Nearest	0.9606